

CAUSALRPC

Distributed computation over Irmin


Craig Ferguson (@CraigFe)

Friday 23rd August




Remote Procedure Call (RPC)

Remote computation at the level of individual function calls.



Trendy microservice architectures



Use RPC for remote access to a resource (storage, private network etc.).

Two birds and one cannonball

Two birds and one cannonball

Problem: consistency under race conditions

- breaks RPC abstraction
- results in additional round-trips to server

Two birds and one cannonball

Problem: consistency under race conditions

- breaks RPC abstraction
- results in additional round-trips to server

Problem: tracing distributed systems

- neither the middleware nor the application layer has full context
- linear log outputs on remote systems must be correlated

Two birds and one cannonball

Problem: consistency under race conditions

- breaks RPC abstraction
- results in additional round-trips to server

Problem: tracing distributed systems

- neither the middleware nor the application layer has full context
- linear log outputs on remote systems must be correlated

Idea: middleware that understands state

- application explains the underlying state to the middleware
- middleware resolves conflicts and constructs a trace as it goes


Introducing Irmin

- ▶ Distributed, immutable Merkle trees with push/pull sync – just like Git!
- ▶ Pure OCaml; unikernel-compatible.
- ▶ Parameterised on:
 - ▶ Storage backend (Memory, FS, Git repo, Redis, ...)
 - ▶ User-defined mergable contents (leaves of the tree)

Aside: why three-way merge?

Aside: why three-way merge?

Two-way CRDT¹ merge




state → state → state

¹M. Shapiro et al. *Conflict-Free Replicated Data Types*, 2011.


Aside: why three-way merge?

Two-way CRDT¹ merge



state → state → state

Three-way Irmin merge

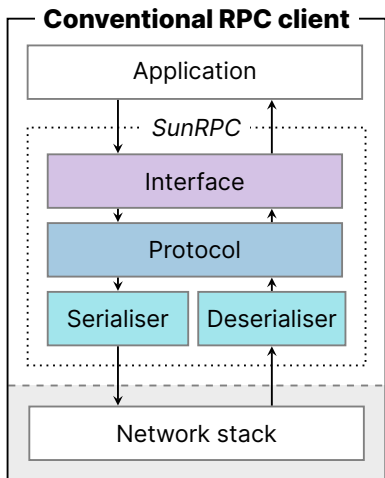


lca : state → state → state → state

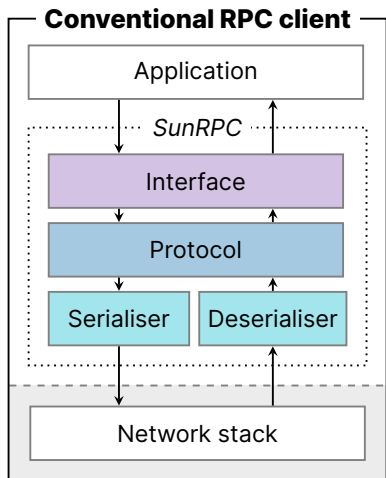
Lowest common ancestor allows the merge operation to reason about *intent*.

¹M. Shapiro et al. *Conflict-Free Replicated Data Types*, 2011.

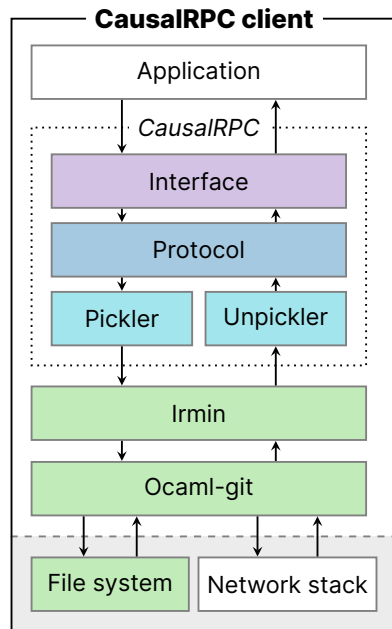
State-aware RPC



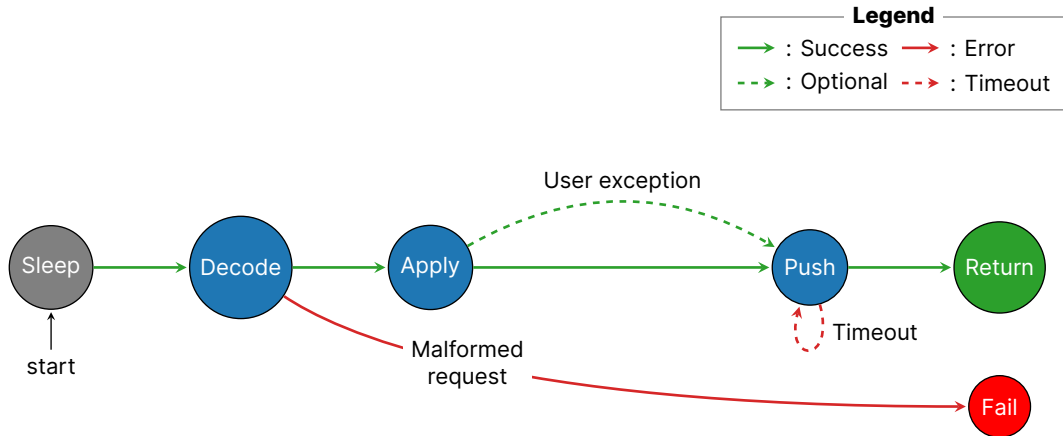
State-aware RPC



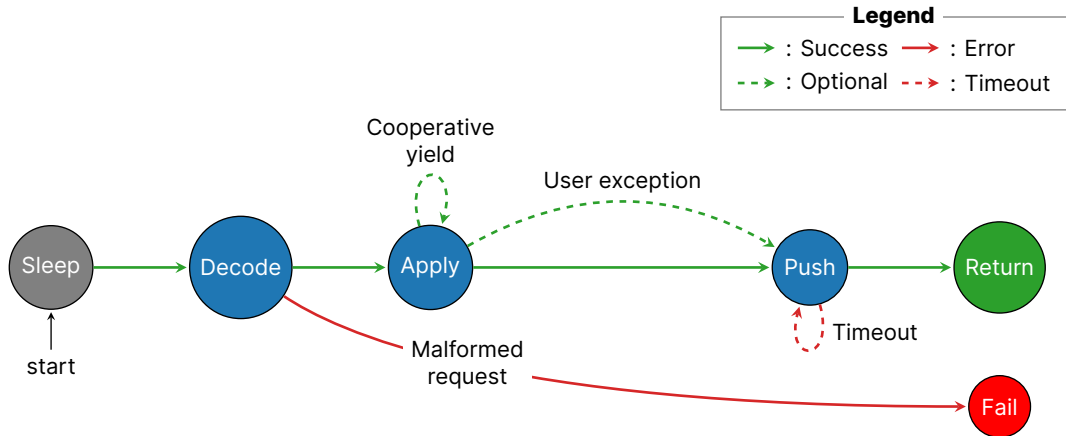
*Expose underlying
statefulness*



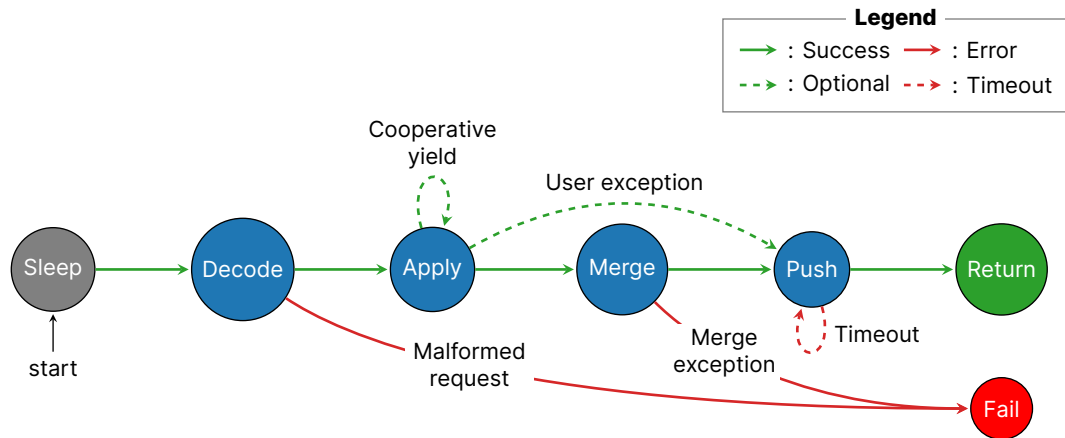
The life of an RPC thread



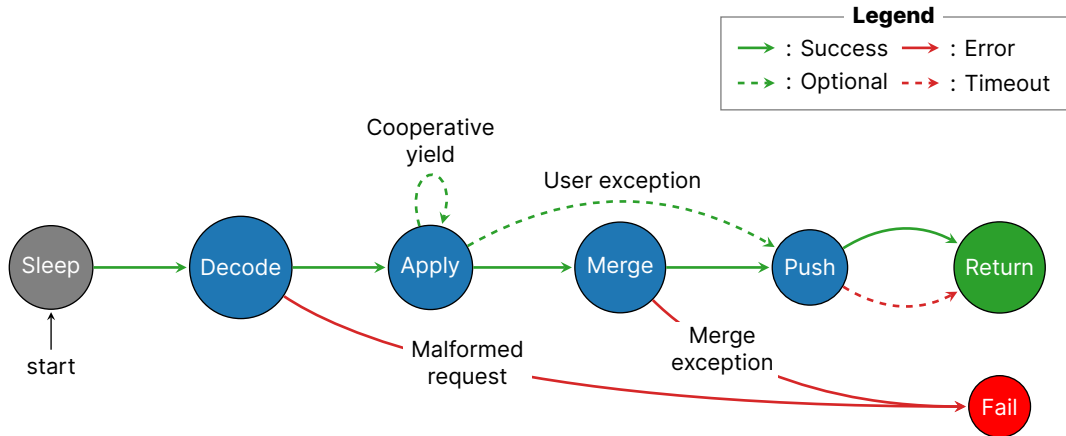
The life of an RPC thread




The life of an RPC thread



The life of an RPC thread



Mergeable filesystems



Mergeable filesystems

client-a

server



client-b

server


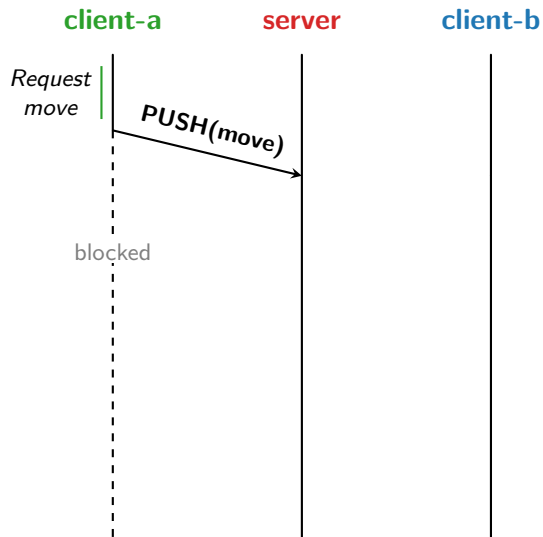


*Initial
world-view*


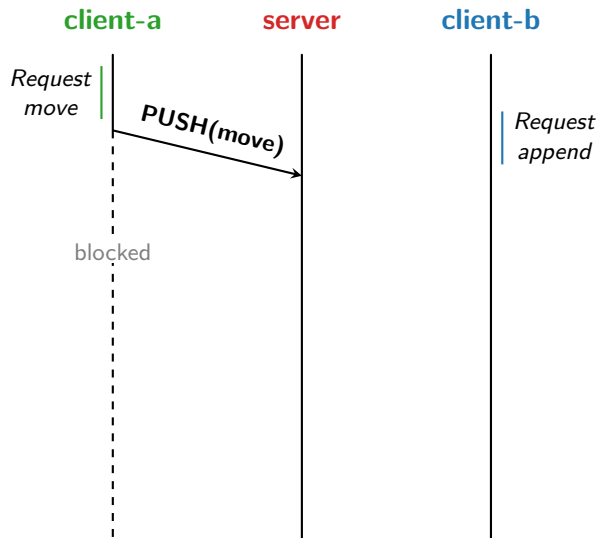
Mergeable filesystems




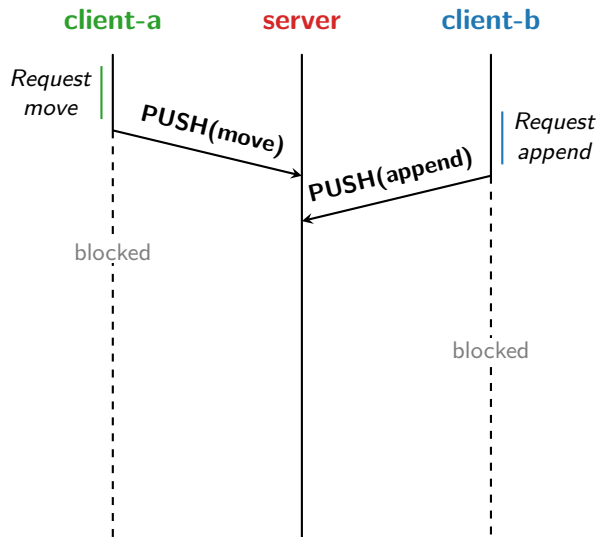
Mergeable filesystems




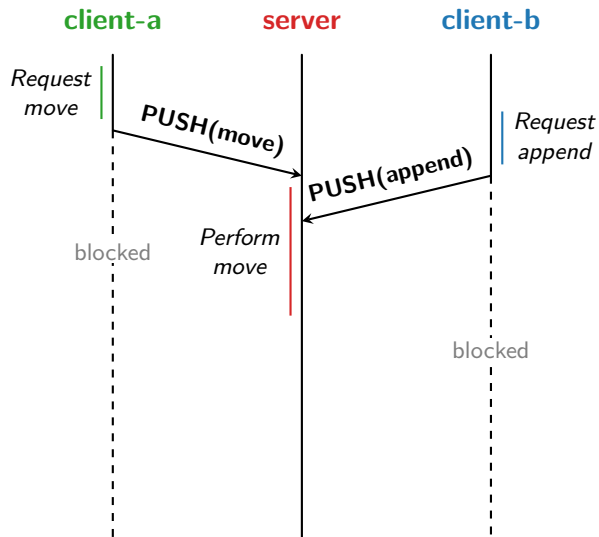
Mergeable filesystems




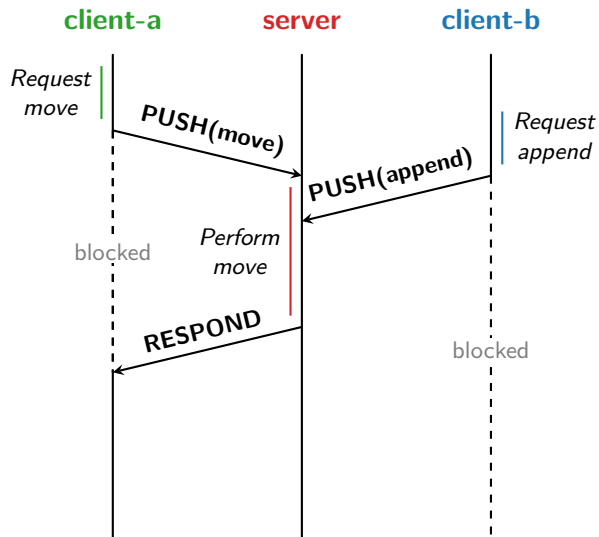
Mergeable filesystems





Mergeable filesystems





Mergeable filesystems




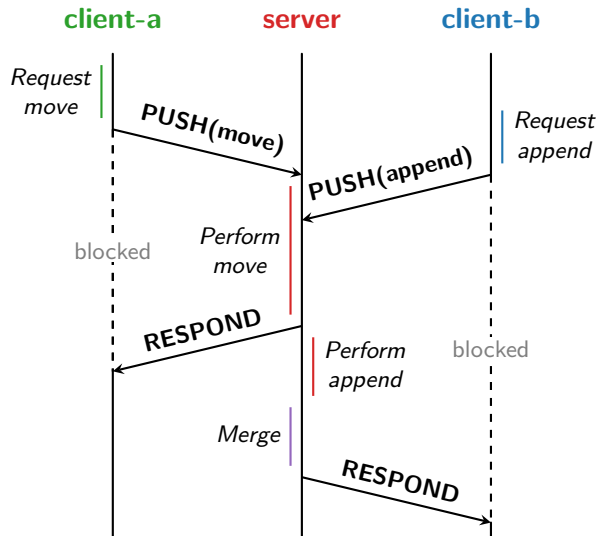
Mergeable filesystems



Mergeable filesystems



Mergeable filesystems



Mergeable filesystems: trace output

```
$ git log
```

Mergeable filesystems: trace output

```
$ git log
```

```
* server Initialised Filesystem store
```

```
|\
```

```
* | clientA Request <move>: /tmp/sharedFile -> /log/sharedFile
```

```
* | server (client--a) Perform <move> for clientA
```

```
| * clientB Request <append>: /tmp/sharedFile
```

```
| * server Perform <append> operation for clientB
```

```
|/
```

```
* server (server, client--b) Merge client--b into server
```

Why OCaml?

Portability – Inherits the Mirage mindset towards platform-independence.

Why OCaml?

Portability – Inherits the Mirage mindset towards platform-independence.

Type-safety – Strict interface definition using an embedded DSL.


Why OCaml?

Portability – Inherits the Mirage mindset towards platform-independence.

Type-safety – Strict interface definition using an embedded DSL.


Coolness-factor (!!!) – Protocol design in OCaml is really fun.

Performance: latency




Measured over 1000 RPCs for a single client-server pair over a Docker bridge network.

Server run-time cost profile from first packet receipt to final packet send:



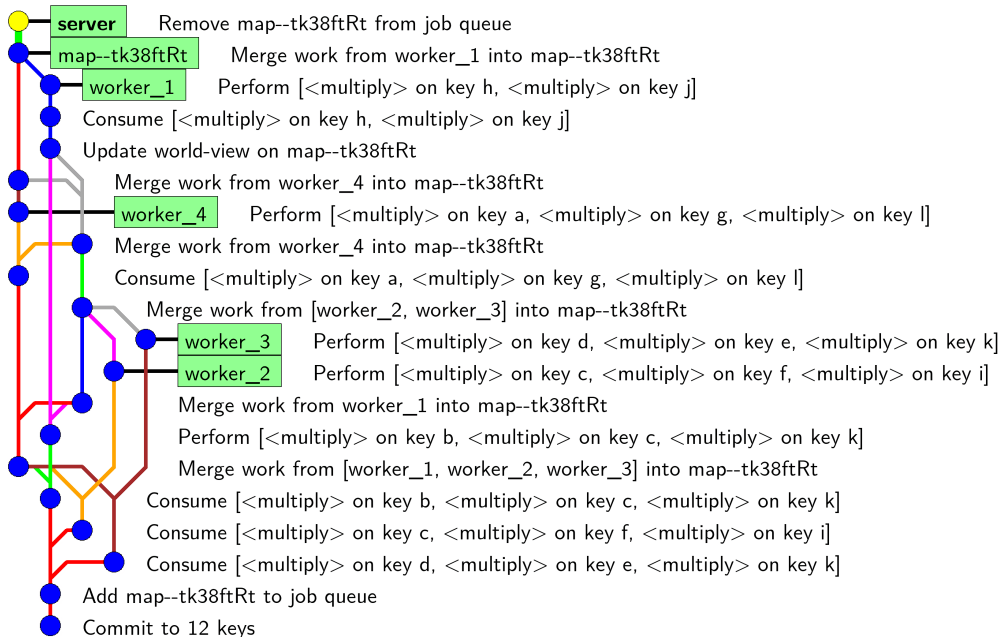
Source: *gprof* and a huge amount of patience.

Performance: memory-usage



Generalisation to work clusters

Generalisation to work clusters




Work cluster scalability


Compare cluster performance to Amdahl's law prediction:

Work cluster scalability


Compare cluster performance to Amdahl's law prediction:




What next?



What next?




What next?




¹N. Krishnaswami and J. Yallop. *A Typed, Algebraic Approach to Parsing*, 2019.

What next?



¹N. Krishnaswami and J. Yallop. *A Typed, Algebraic Approach to Parsing*, 2019.

What next?



¹N. Krishnaswami and J. Yallop. *A Typed, Algebraic Approach to Parsing*, 2019.

²A. Mokhov et al. *Selective Applicative Functors*, 2019.

Take-away lessons

1. Simple distributed system abstractions set challenging design constraints.

Take-away lessons

1. Simple distributed system abstractions set challenging design constraints.
2. Improve semantics by squashing the stack; use OCaml to maintain abstraction and flexibility.

Take-away lessons

1. Simple distributed system abstractions set challenging design constraints.
2. Improve semantics by squashing the stack; use OCaml to maintain abstraction and flexibility.
3. Consider non-linear log structures in your programs; *casual* logs make it easy to reason about intent.

**Thanks For
Listening**

Related work

(Much) more detail in the full dissertation:

<https://craigfe.io/out/causalrpc.pdf>

mergeable types

- ▶ Conflict-free replicated datatypes. [M. Shapiro et al. 2011]
- ▶ Mergeable types. [B Farinier et al. 2015]

embedded DSLs for RPC

- ▶ Pickler combinators. [A. Kennedy. 2004]

patterns for remote execution

- ▶ Remote functors. [A. Gill et al. 2015]
- ▶ Selective applicative functors. [A. Mokhov et al. 2019]

Work cluster performance

